



Building Cloud Native Applications on Cloud Foundry

An in depth look at the microservices architecture pattern,
containers and Cloud Foundry

10: CFCR/PKS and Kubernetes

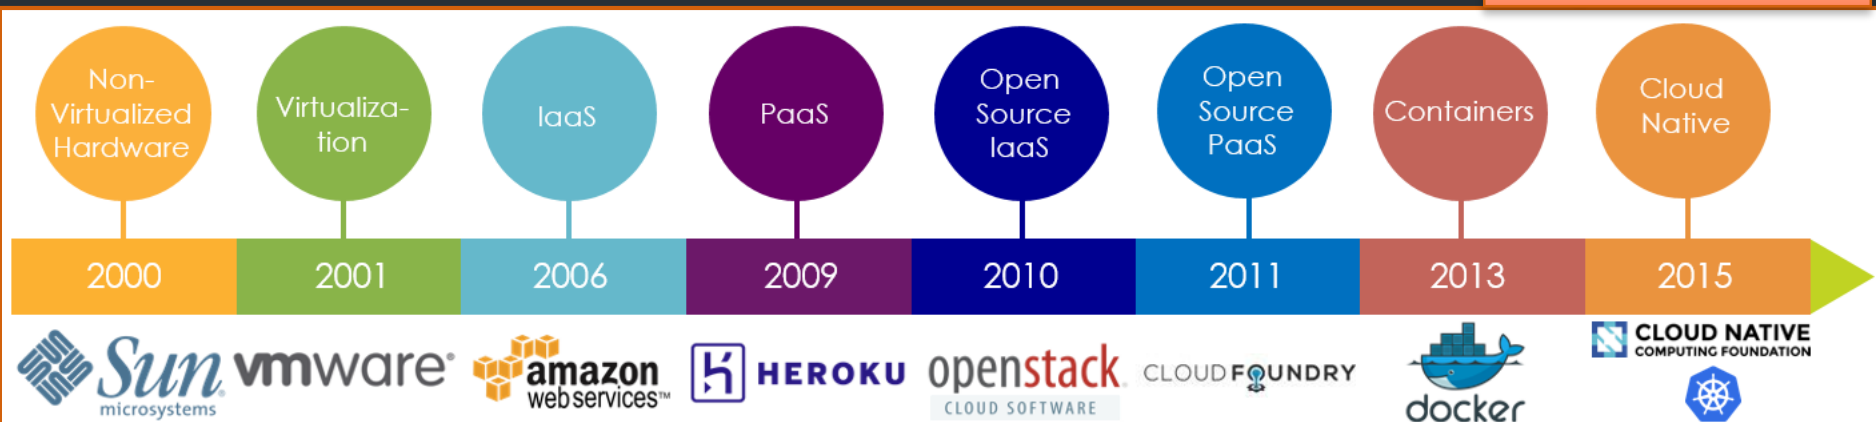
Objectives

- Describe the goals of the Kubernetes system
- Examine the design principles of Kubernetes
- Explore the Kubernetes architecture
- List the Kubernetes service components

Platform Evolution

- The platform underlying rapidly growing Internet applications has changed significantly over the years
 - .com era – bare metal servers
 - Web 2.0 – IaaS
 - APIs – PaaS
 - Microservices – Cloud Native

Image courtesy CNCF



Important PaaS

5

Copyright 2017-2018, RX-M LLC

- Public Cloud based

- AppEngine

- Because it is Google, first major PaaS offering

- Heroku/SalesForce

- The most popular web tech PaaS (Ruby on Rails, Node.js Express, Python Django)

- Elastic Beanstalk

- Because it is Amazon and linked to the services of the largest cloud

- Azure

- Because it is Microsoft and the cornerstone of .net in the cloud

- Open Source PaaS

- Cloud Foundry

- The Java Spring community centric PaaS

- OpenShift

- The Java EE community centric PaaS



The Cloud Foundry Foundation

6

Copyright 2017 - 2018, RX-M LLC

PLATINUM



GOLD



SILVER

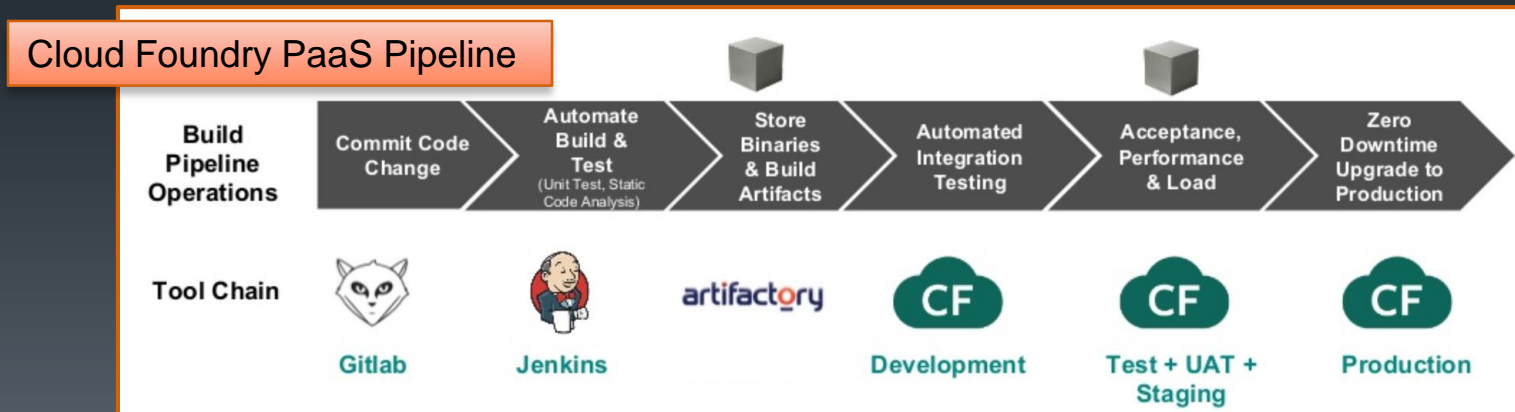


PaaS - CFAR

7

Copyright 2017 - 2018, RX-M LLC

- Platform as a Service
 - Push code
 - PaaS supports only applications fitting one of it's **buildpack** models
 - The PaaS builds, tests, packages and deploys
 - Platform containerizes the artifacts
 - Perhaps JARs, a custom format, Docker Images or other
- Cloud Foundry Application Runtime [CFAR] is the Cloud Foundry PaaS



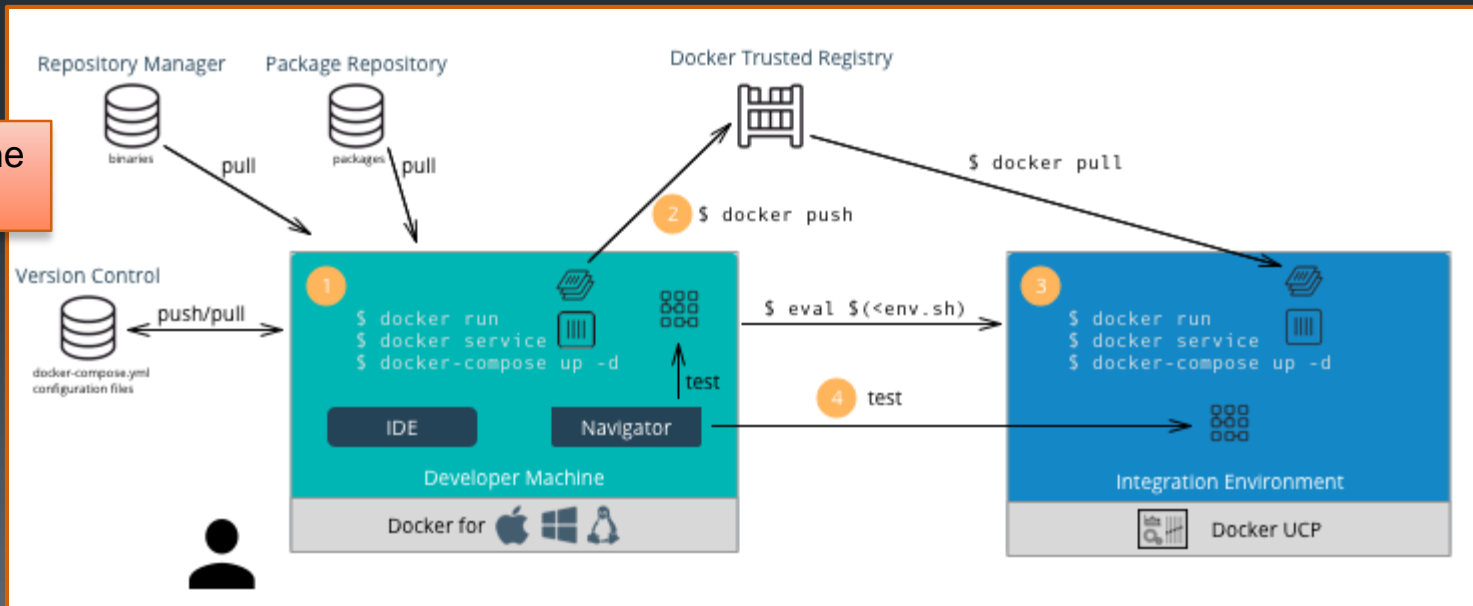
CaaS - CFCR

8

Copyright 2017 - 2018, RX-M LLC

- Containers as a Service
 - Push container Images
 - The CaaS deploys and orchestrates
 - Containers are created earlier in the pipeline reducing variability
- Hybrid
 - Most popular PaaS today use container technology
 - Cloud Foundry uses containers and its own OCI runc based Guardian container manager (Greenhouse on Windows)
 - RedHat OpenShift uses Docker as the container manager and Kubernetes as the orchestration engine
- Cloud Foundry Container Runtime [CFCR] is Cloud Foundry's Kubernetes solution

Docker Dev Pipeline Reference Arch

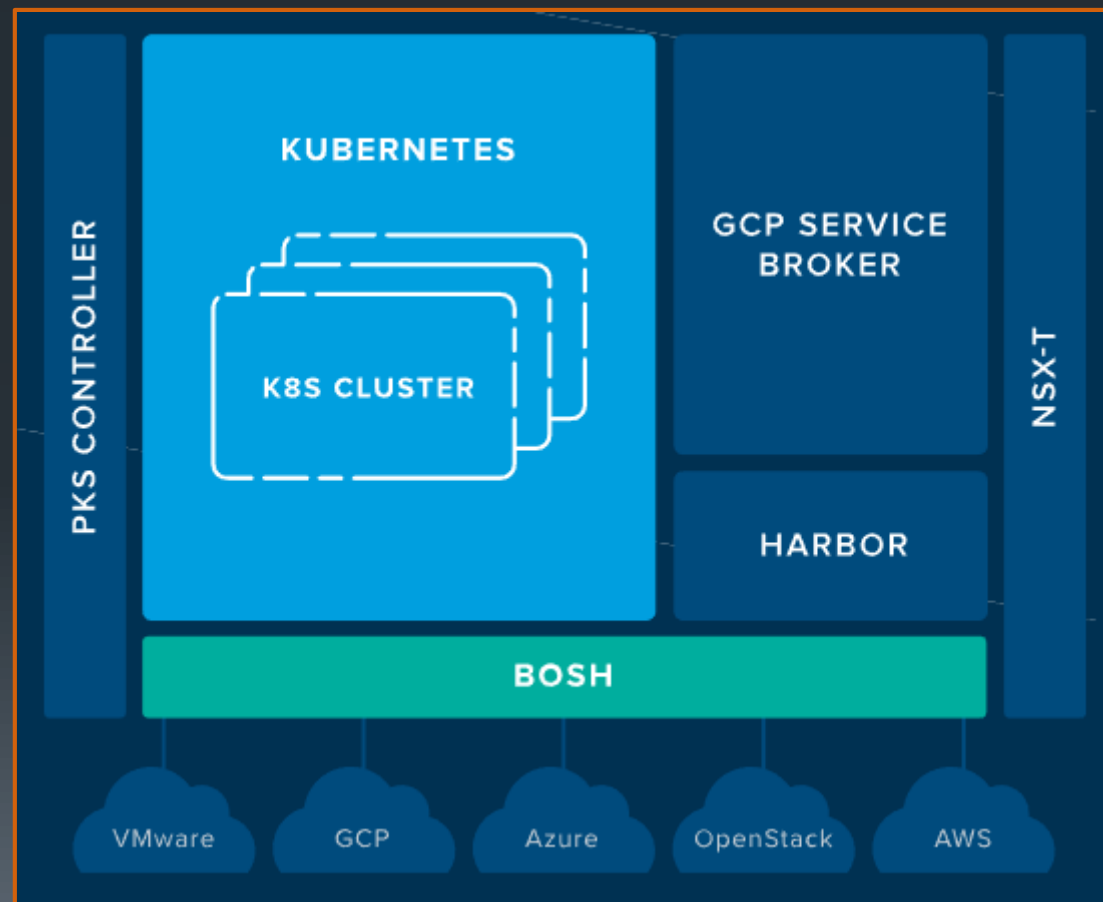


Pivotal Container Service [PKS]


9

Copyright 2017 - 2018, RX-M LLC

- **Kubernetes**
 - Current stable release of Kubernetes, deployed and managed by BOSH
 - No proprietary extensions
- **PKS Controller**
 - Creates and scales Kubernetes clusters from the command line and API
- **GCP Service Broker**
 - Allows apps to access Google Cloud APIs
 - Provides Container Engine (GKE) portability
- **Harbor**
 - A docker registry server with vulnerability scanning, identity management, and support for multiple registries
- **NSX-T**
 - VMware NSX-T secure SDN
- **BOSH**
 - BOSH deploys, scales, monitors and heals Kubernetes clusters



Google Trends



Kubernetes Design Overview

11

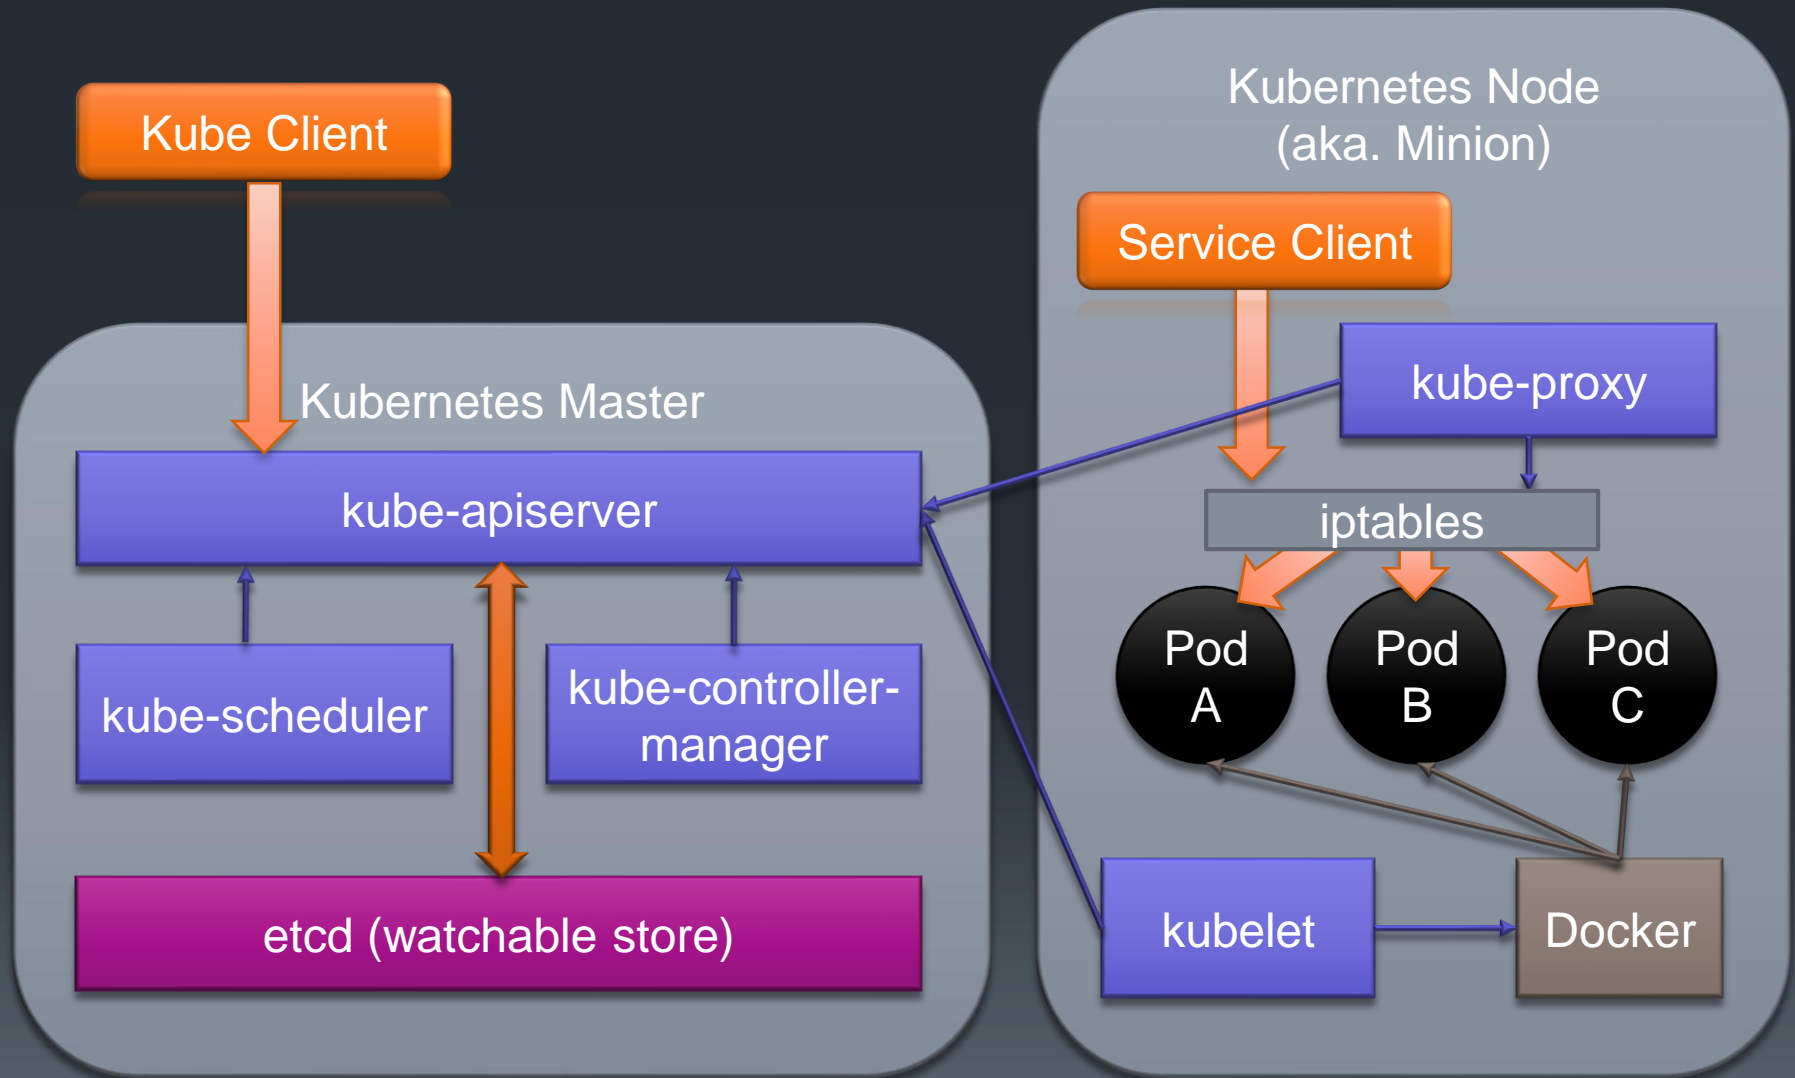
Copyright 2013-2018, RX-M LLC

- Kubernetes is a system for managing containerized applications across multiple hosts
 - Deployment
 - Maintenance
 - Scaling
- Kubernetes establishes robust declarative primitives for maintaining the desired state requested by the user
 - Main value add of Kubernetes, **declarative target state**
- Kubernetes is primarily targeted at applications composed of multiple containers
 - Elastic distributed **micro-services**
- Also designed to facilitate **migration of non-containerized application stacks** to Kubernetes
 - Includes abstractions for grouping containers in both loosely coupled and tightly coupled formations
 - Provides ways for containers to find and communicate with each other in relatively familiar ways
- Kubernetes enables users to ask a cluster to run a set of containers
 - The system automatically chooses hosts to run those containers on
 - Scheduling is a policy-rich, topology-aware, workload-specific function that significantly impacts:
 - Availability
 - Performance
 - Capacity
 - The scheduler takes into account:
 - individual and collective resource requirements
 - quality of service requirements
 - hardware/software/policy constraints
 - affinity and anti-affinity specifications
 - data locality
 - inter-workload interference
 - Deadlines
 - Workload-specific requirements exposed through the API as necessary
- **Self-healing mechanisms**
 - Auto-restart
 - Re-schedule
 - Replication
- Kubernetes **runs on a number of cloud providers**, as well as on physical hosts
- A single Kubernetes cluster is **not intended to span multiple availability zones**
 - Recommend building a higher-level layer to replicate complete deployments of highly available applications across multiple zones
- Kubernetes aspires to be an **extensible, pluggable**, building-block OSS platform and toolkit
 - Kubernetes is built as a collection of pluggable components and layers
 - The ability to use alternative schedulers, controllers, storage systems, and distribution mechanisms
 - Others are building higher-level PaaS functionality and multi-cluster layers, without modification of core Kubernetes source
 - API isn't just (or even necessarily mainly) targeted at end users, but at tool and extension developers
 - All APIs are visible and available, including the APIs used by the scheduler, the node controller, the replication-controller manager, etc.

Kubernetes Architecture

12

Copyright 2013-2018, RX-M LLC



Master Components

13

Copyright 2013-2018, RX-M LLC

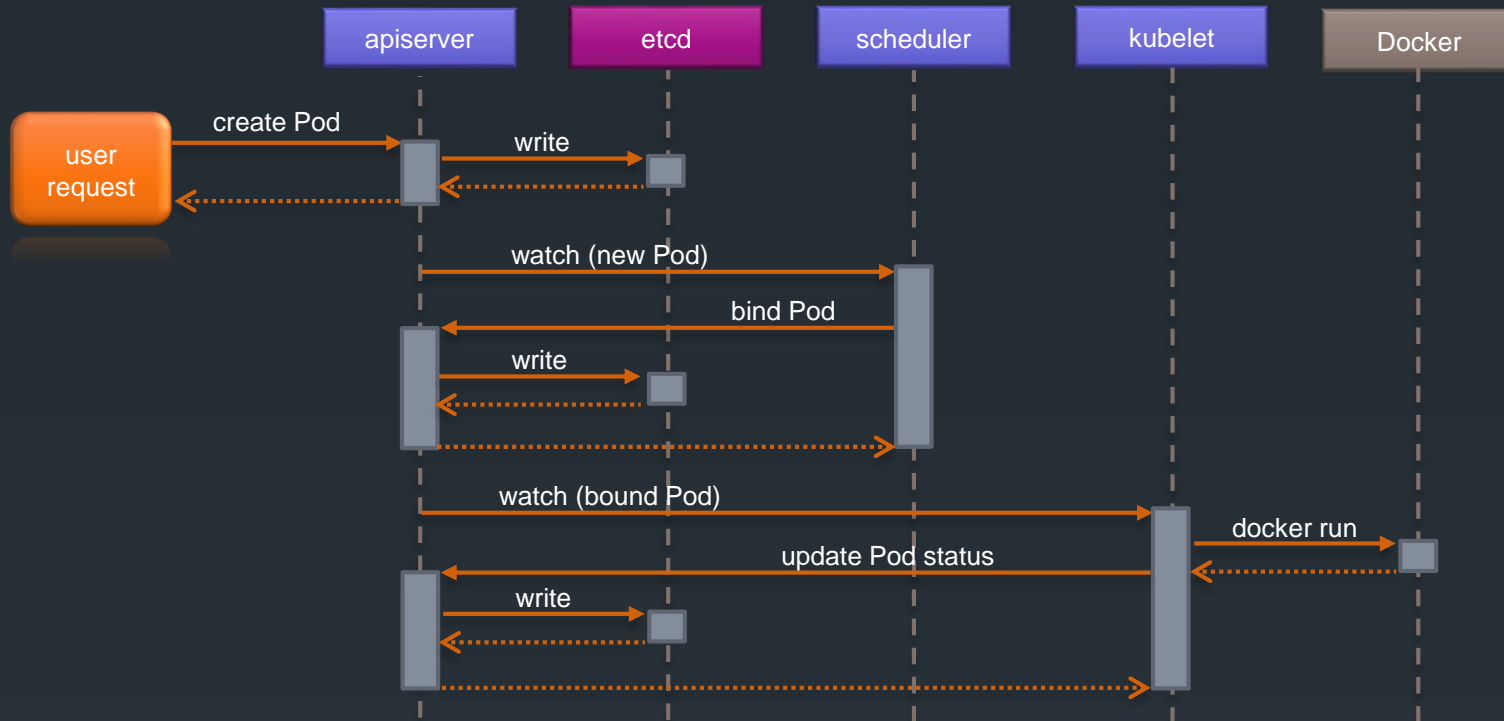
- The Kubernetes Control Plane
 - The Kubernetes control plane is split into a set of components which currently all run on a master node
 - These components work together to provide a unified view of the cluster
- **etcd**
 - All persistent master state is stored in an instance of etcd
 - This provides a way to store configuration data reliably
 - Watch support allows coordinating components to be notified of changes
- **API Server**
 - The apiserver serves the Kubernetes API
 - A CRUD-y server, with most/all business logic implemented in separate components or in plug-ins
 - Mainly processes REST operations, validates them, and updates the corresponding objects in etcd
- **Scheduler**
 - Binds unscheduled pods to nodes via the /binding API
 - The scheduler is pluggable
- **Controller Manager**
 - All other cluster-level functions are currently performed by the Controller Manager
 - Endpoint objects are created and updated by the endpoints controller
 - Nodes are discovered, managed, and monitored by the node controller
 - Replication is managed by Deployments & Replica Sets
 - Could eventually be split into separate components to make them independently pluggable

- Every resource in Kubernetes, such as a pod, is identified by a URI and has a UID
- The URI includes
 - kind of object (e.g. pod)
 - object name
 - Namespace
- For a certain object kind, every name is unique within its namespace
- In contexts where an object name is provided without a namespace, it is assumed to be in the default namespace

Pod Creation Flow

14

Copyright 2013-2018, RX-M LLC



- The user creates a Pod via the API Server and the API server writes it to etcd
- The scheduler notices an “unbound” Pod and decides which node to run that Pod on
 - It writes that binding back to the API Server
- The Kubelet notices a change in the set of Pods that are bound to its node
 - It, in turn, runs the container via the container runtime (Docker)
- The Kubelet monitors the status of the Pod via the container runtime
 - As things change, the Kubelet will reflect the current status back to the API Server

■ Service listing from a standalone Kubernetes cluster:

```
user@ubuntu:~$ sudo ps -ef | grep kube
```

```
kube-apiserver --client-ca-file=/etc/kubernetes/pki/ca.crt --tls-cert-file=/etc/kubernetes/pki/apiserver.crt
--tls-private-key-file=/etc/kubernetes/pki/apiserver.key --proxy-client-cert-file=/etc/kubernetes/pki/front-proxy-client.crt
--service-account-key-file=/etc/kubernetes/pki/sa.pub --kubelet-client-key=/etc/kubernetes/pki/apiserver-kubelet-client.key
--proxy-client-key-file=/etc/kubernetes/pki/front-proxy-client.key --requestheader-group-headers=X-Remote-Group
--requestheader-extra-headers-prefix=X-Remote-Extra- --requestheader-allowed-names=front-proxy-client --secure-port=6443
--kubelet-client-certificate=/etc/kubernetes/pki/apiserver-kubelet-client.crt
--requestheader-client-ca-file=/etc/kubernetes/pki/front-proxy-ca.crt --insecure-port=0 --allow-privileged=true
--experimental-bootstrap-token-auth=true --kubelet-preferred-address-types=InternalIP,ExternalIP,Hostname
--requestheader-username-headers=X-Remote-User --service-cluster-ip-range=10.96.0.0/12 --authorization-mode=Node,RBAC
--advertise-address=192.168.225.179 --etcd-servers=http://127.0.0.1:2379 --admission-control=Initializers,NamespaceLifecycle,LimitRanger,
ServiceAccount,PersistentVolumeLabel,DefaultStorageClass,DefaultTolerationSeconds,NodeRestriction,ResourceQuota

kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfig=/etc/kubernetes/kubelet.conf
--pod-manifest-path=/etc/kubernetes/manifests --allow-privileged=true --network-plugin=cni --cni-conf-dir=/etc/cni/net.d
--cni-bin-dir=/opt/cni/bin --cluster-dns=10.96.0.10 --cluster-domain=cluster.local --authorization-mode=Webhook
--client-ca-file=/etc/kubernetes/pki/ca.crt --cadvisor-port=0 --rotate-certificates=true --cert-dir=/var/lib/kubelet/pki

kube-controller-manager --address=127.0.0.1 --leader-elect=true --controllers=*,bootstrapsigner,tokencleaner --root-ca-file=/etc/kubernetes/pki/ca.crt
--service-account-private-key-file=/etc/kubernetes/pki/sa.key --cluster-signing-key-file=/etc/kubernetes/pki/ca.key
--use-service-account-credentials=true --kubeconfig=/etc/kubernetes/controller-manager.conf
--cluster-signing-cert-file=/etc/kubernetes/pki/ca.crt

kube-scheduler --leader-elect=true --kubeconfig=/etc/kubernetes/scheduler.conf --address=127.0.0.1

kube-proxy --kubeconfig=/var/lib/kube-proxy/kubeconfig.conf

kube-dns --domain=cluster.local. --dns-port=10053 --config-dir=/kube-dns-config --v=2

sidecar --v=2 --logtostderr --probe=kubedns,127.0.0.1:10053,kubernetes.default.svc.cluster.local,5,A
--probe=dnsmaq,127.0.0.1:53,kubernetes.default.svc.cluster.local,5,A
```

```
user@ubuntu:~$ sudo netstat -ntlp
```

Active Internet connections (only servers)

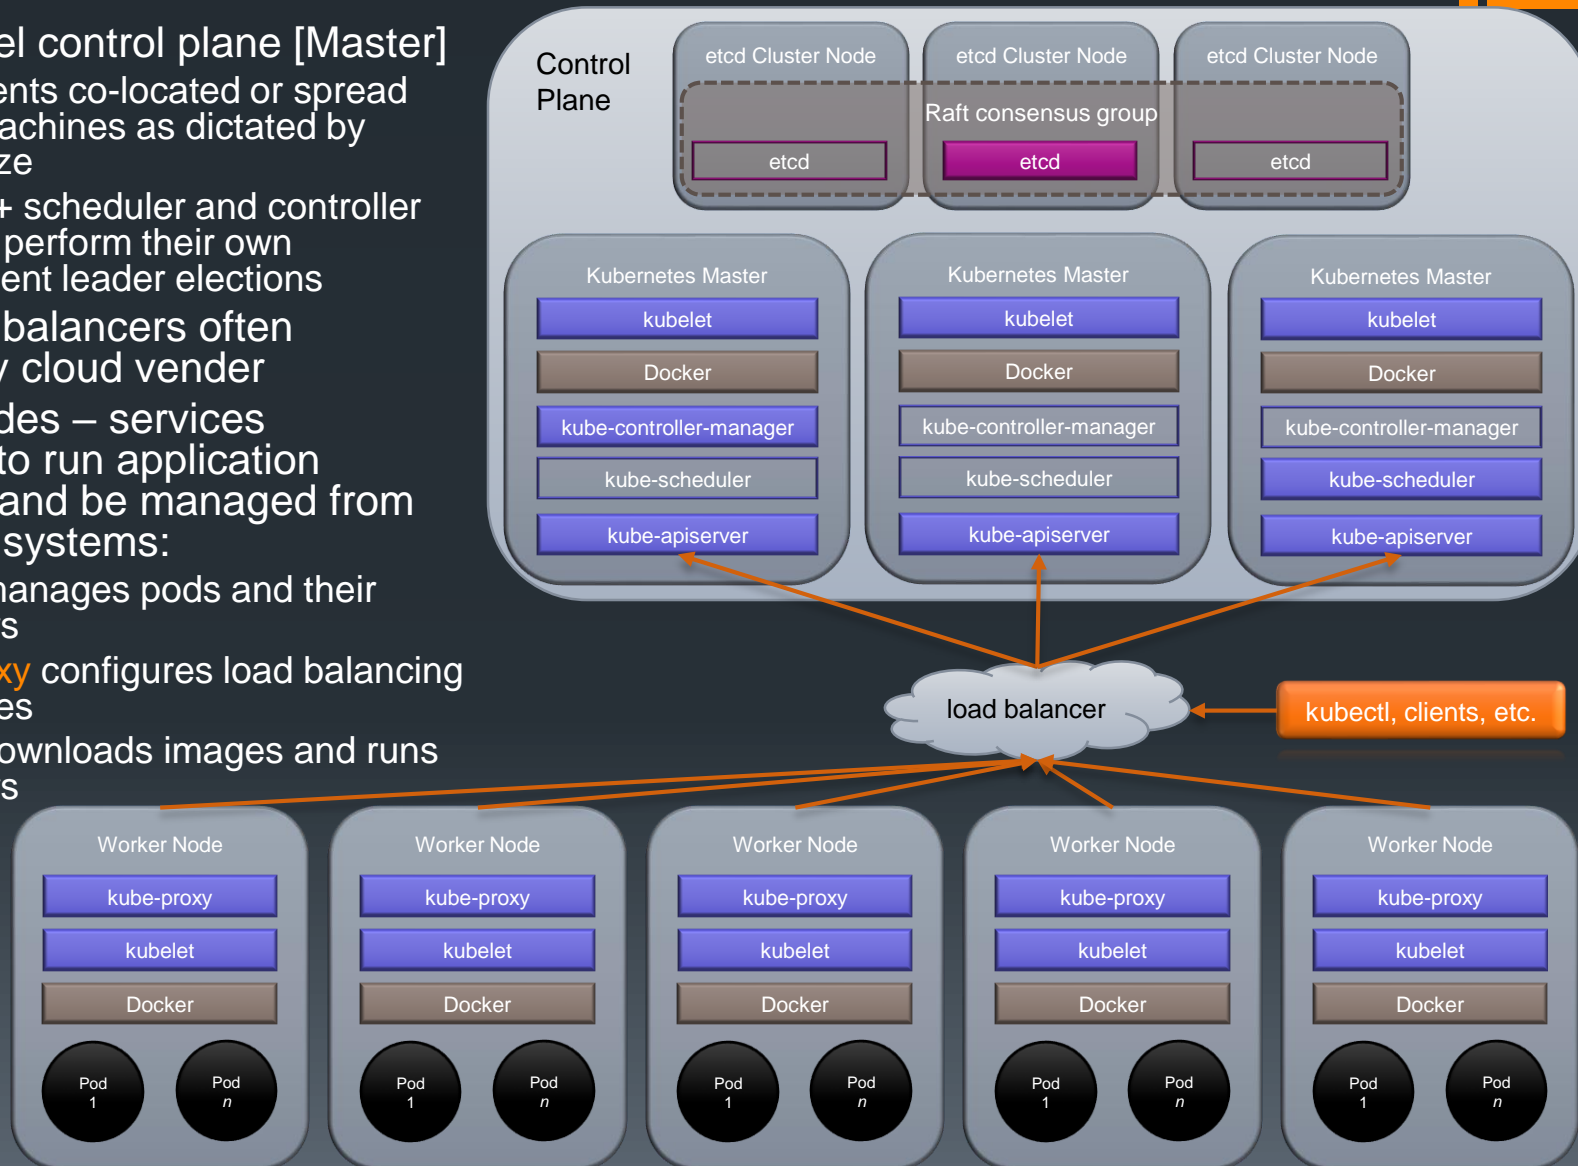
Proto	Recv-Q	Send-Q	Local Address	Foreign Address	State	PID/Program name
tcp	0	0	127.0.0.1:10248	0.0.0.0:*	LISTEN	1058/kubelet
tcp	0	0	127.0.0.1:10249	0.0.0.0:*	LISTEN	2608/kube-proxy
tcp	0	0	127.0.0.1:10251	0.0.0.0:*	LISTEN	1907/kube-scheduler
tcp	0	0	127.0.0.1:10252	0.0.0.0:*	LISTEN	2105/kube-controller
tcp	0	0	127.0.0.1:2379	0.0.0.0:*	LISTEN	2077/etcd
tcp	0	0	127.0.0.1:2380	0.0.0.0:*	LISTEN	2077/etcd
tcp6	0	0	:::10250	:::*	LISTEN	1058/kubelet
tcp6	0	0	:::6443	:::*	LISTEN	2694/kube-apiserver
tcp6	0	0	:::10255	:::*	LISTEN	1058/kubelet
tcp6	0	0	:::10256	:::*	LISTEN	2608/kube-proxy

Scaled Architecture & HA Model

16

Copyright 2013-2018, RX-M LLC

- Distributed watchable storage via etcd & Raft
- Cluster-level control plane [Master]
 - Components co-located or spread across machines as dictated by cluster size
 - K8s v1.2+ scheduler and controller manager perform their own independent leader elections
- Proxy load balancers often provided by cloud vendor
- Worker Nodes – services necessary to run application containers and be managed from the master systems:
 - **kubelet** manages pods and their containers
 - **kube-proxy** configures load balancing via iptables
 - **Docker** downloads images and runs containers



etcd

- OpenSource distributed consistent key-value store for shared configuration and service discovery from CoreOS
 - Originally created as a distributed configuration information tool for CoreOS Linux clusters
 - Placed all of the configuration information for cluster nodes into in the key-value store that could be rapidly searched
 - Named for the /etc directory in Linux (which stores all of the config info about a server) and "d" for distributed systems
- etcd focuses on :
 - Simple**: curl'able user-facing API (v3 uses gRPC)
 - Secure**: automatic TLS w/ optional client cert auth
 - Fast**: benchmarked 1000s of writes/s per instance
 - Resilient**: properly distributed using Raft
- etcd is written in Go
- etcd uses the Raft consensus algorithm to manage a highly-available replicated log
 - etcd can recover from hardware failure and network partitions
- etcdctl** is a simple command line client that can be used in scripts or to explore an etcd cluster
- Kubernetes** uses etcd for storing and replicating data across the entire cluster
 - Cloud Foundry** also uses etcd as their distributed key-value store
 - Docker** supports etcd for and multi-host networking
- Used to track pod deployment, execution status, labels, endpoints, etc.

Over 500 projects on GitHub are making use of etcd in one form or another: <https://github.com/coreos/etcd>

```
user@ubuntu:~$ etcdctl ls --recursive | grep kube-system
/registry/services/endpoints/kube-system
/registry/services/endpoints/kube-system/kube-controller-manager
/registry/services/endpoints/kube-system/kube-dns
/registry/services/endpoints/kube-system/kube-scheduler
/registry/services/endpoints/kube-system/kubernetes-dashboard
/registry/services/specs/kube-system
/registry/services/specs/kube-system/kube-dns
```

Service Name and Transport Protocol Port Number Registry

Last Updated
2016-01-13

Expert(s)

TCP/UDP: Joe Touch, Eliot Lear, Allison Mankin, Markku Kojo, Kumiko Ono, Martin Stiermerling,
Lars Eggert, Alexey Melnikov, Wes Eddy, and Alexander Zimmermann
SCTP: Allison Mankin and Michael Tuexen
DCCP: Eddie Kohler and Yoshifumi Nishida

Reference

[RFC6335]

Note

Service names and port numbers are used to distinguish between different services that run over transport protocols such as TCP, UDP, DCCP, and SCTP.

Service names are assigned on a first-come, first-served process, as documented in [RFC6335].

Port numbers are assigned in various ways, based on three ranges: System Ports (0-1023), User Ports (1024-49151), and the Dynamic and/or Private Ports (49152-65535); the difference uses of these ranges is described in [RFC6335]. System Ports are assigned by IETF process for standards-track protocols, as per [RFC6335]. User Ports are assigned by IANA using the "IETF Review" process, the "IESG Approval" process, or the "Expert Review" process, as per [RFC6335]. Dynamic Ports are not assigned.

The registration procedures for service names and port numbers are described in [RFC6335].

Assigned ports both System and User ports SHOULD NOT be used without or prior to IANA registration.

```
*****
* PLEASE NOTE THE FOLLOWING:                                     *
* * * * *
* ASSIGNMENT OF A PORT NUMBER DOES NOT IN ANY WAY IMPLY AN      *
* ENDORSEMENT OF AN APPLICATION OR PRODUCT, AND THE FACT THAT NETWORK *
* TRAFFIC IS FLOWING TO OR FROM A REGISTERED PORT DOES NOT MEAN THAT *
* IT IS "GOOD" TRAFFIC, NOR THAT IT NECESSARILY CORRESPONDS TO THE *
* ASSIGNED SERVICE. FIREWALL AND SYSTEM ADMINISTRATORS SHOULD *
* CHOOSE HOW TO CONFIGURE THEIR SYSTEMS BASED ON THEIR KNOWLEDGE OF *
* THE TRAFFIC IN QUESTION, NOT WHETHER THERE IS A PORT NUMBER *
* REGISTERED OR NOT.                                           *
*****
```

Available Formats

CSV XML HTML Plain text

etcd Search

Service Name	Port Number	Transport Protocol	Description	Assignee	Contact	Registration Date
etcd-client	2379	tcp	etcd client communication	[CoreOS]	[Brian_Harrington]	2014-07-09
etcd-server	2380	tcp	etcd server to server communication	[CoreOS]	[Brian_Harrington]	2014-07-09

```
/registry/pods/kube-system/kubernetes-dashboard-525851760-62cgm
/registry/pods/kube-system/etcd-ubuntu
/registry/pods/kube-system/kube-apiserver-ubuntu
/registry/pods/kube-system/kube-discovery-1769846148-c3cg4
/registry/serviceaccounts/kube-system
/registry/serviceaccounts/kube-system/default
/registry/deployments/kube-system
/registry/deployments/kube-system/kube-discovery
/registry/deployments/kube-system/kube-dns
/registry/deployments/kube-system/kubernetes-dashboard
```



25081d

Google Kubernetes Engine Example

Part 1 [docker/gcloud]

```
user@ubuntu:~$ mkdir kubeex
user@ubuntu:~$ cd kubeex/
user@ubuntu:~/kubeex$ vim server.js
user@ubuntu:~/kubeex$ cat server.js
var http = require('http');
```

```
var handleRequest = function(request, response) {
  response.writeHead(200);
  response.end("Hello World!");
}
```

```
var www = http.createServer(handleRequest);
www.listen(8080);
```

```
user@ubuntu:~/kubeex$ vim Dockerfile
```

```
user@ubuntu:~/kubeex$ cat Dockerfile
```

```
FROM node:latest
```

```
COPY server.js /server.js
```

```
CMD node server.js
```

```
user@ubuntu:~/kubeex$ docker image build -t gcr.io/kubeex-160118/hello-node .
```

```
Sending build context to Docker daemon 3.072 kB
```

```
Sending build context to Docker daemon
```

```
Step 0 : FROM node:latest
```

```
...
```

```
---> 7b7b1aed5fed
```

```
Step 1 : COPY server.js /server.js
```

```
...
```

```
---> 578bc181cf79
```

```
Step 2 : CMD node server.js
```

```
...
```

```
---> b75c5d346af0
```

```
Successfully built b75c5d346af0
```

```
$ gcloud docker -- image push gcr.io/kubeex-160118/hello-node
```

```
The push refers to a repository [gcr.io/kubeex-160118/hello-node]
```

```
5194b882e159: Pushed
```

```
06ddfcaca7da: Layer already exists
```

```
...
```

```
a2ae92ffcd29: Mounted from google_containers/hyperkube-amd64
```

```
latest: digest: sha256:4c5722d9360a4b0450b46df06dc50fcf5e137944c10a38d9f4b84c88ca88e877 size: 2003
```

```
$ gcloud container clusters create hello-world --num-nodes 1 --machine-type g1-small
```

```
Creating cluster hello-world...done.
```

```
Created [https://container.googleapis.com/v1/projects/kubeex-160118/zones/us-west1-b/clusters/hello-world].
```

```
kubeconfig entry generated for hello-world.
```

NAME	ZONE	MASTER_VERSION	MASTER_IP	MACHINE_TYPE	NODE_VERSION	NUM_NODES	STATUS
hello-world	us-west1-b	1.7.6	104.196.246.216	g1-small	1.7.6	1	RUNNING

```
$ gcloud compute instances list
```

NAME	ZONE	MACHINE_TYPE	PREEMPTIBLE	INTERNAL_IP	EXTERNAL_IP	STATUS
gke-hello-world-default-pool-e7535bd3-pskk	us-west1-b	g1-small		10.138.0.2	35.185.199.37	RUNNING

N.B. The gcloud command must be installed and configured to use a valid Google Cloud account

1. Create an application container
2. Push the container to the Google Container Registry
3. Create a Google Compute cluster to run containers on

Google Kubernetes Engine Example Part 2 [kubectl]

```
$ kubectl run hello-node --image=gcr.io/kubeex-160118/hello-node --port=8080
```

```
deployment "hello-node" created
```

```
$ kubectl get deploy
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
hello-node	1	1	1	0	14s

```
$ kubectl expose deployment hello-node --type="LoadBalancer"
```

```
service "hello-node" exposed
```

```
$ kubectl get services hello-node
```

NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
hello-node	10.3.254.37	35.185.198.254	8080:31487/TCP	1m

```
$ kubectl scale deployment hello-node --replicas=3
```

```
deployment "hello-node" scaled
```

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
hello-node-952813049-2pf00	1/1	Running	0	49s
hello-node-952813049-dccn2	1/1	Running	0	49s
hello-node-952813049-zkxgf	1/1	Running	0	49s

```
$ kubectl describe service hello-node
```

```
Name: hello-node
Namespace: default
Labels: run=hello-node
Selector: run=hello-node
Type: LoadBalancer
IP: 10.3.254.37
LoadBalancer Ingress: 35.185.198.254
Port: <unset> 8080/TCP
NodePort: <unset> 31487/TCP
Endpoints: 10.0.0.11:8080,10.0.0.10:8080,10.0.0.9:8080
Session Affinity: None
Events:
```

FirstSeen	LastSeen	Count	From	SubObjectPath	Type	Reason	Message
-----	-----	-----	----	-----	-----	-----	-----
12m	12m	1	{service-controller }		Normal	CreatingLoadBalancer	Creating load balancer
11m	11m	1	{service-controller }		Normal	CreatedLoadBalancer	Created load balancer

```
$ curl http://35.185.198.254:8080
```

```
Hello World!
```

1. Run a one container pod
2. Expose the pod as a load balanced service
3. Dynamically scale up the pod count in the service

```
$ kubectl describe pod hello-node-952813049-zkxgf
Name:                hello-node-952813049-zkxgf
Namespace:           default
Node:                gke-hello-world-default-pool-e7535bd3-pskk/10.138.0.2
Start Time:          Tue, 28 Feb 2017 14:23:45 -0800
Labels:              pod-template-hash=952813049
                    run=hello-node
Status:              Running
IP:                  10.0.0.10
Controllers:         ReplicaSet/hello-node-952813049
Containers:
  hello-node:
    Container ID:     docker://b0b3812eb6dc2cf867452d5bc49d3345e641c70f3ae75048b180b6d7bed02ec7
    Image:             gcr.io/kubeex-160118/hello-node
    Image ID:          docker://sha256:61dc0720f1441331b6a368dae57f3ab79ecc485e992521fb4268fabbc901d7c1
    Port:              8080/TCP
    Requests:
      cpu:              100m
    State:              Running
      Started:          Tue, 28 Feb 2017 14:23:46 -0800
    Ready:              True
    Restart Count:      0
    Volume Mounts:      /var/run/secrets/kubernetes.io/serviceaccount from default-token-5qbk1 (ro)
    Environment Variables:
      <none>
Conditions:
  Type              Status
  Initialized        True
  Ready              True
  PodScheduled       True
Volumes:
  default-token-5qbk1:
    Type:              Secret (a volume populated by a Secret)
    SecretName:         default-token-5qbk1
QoS Class:           Burstable
Tolerations:         <none>
Events:
...
```

```
$ kubectl delete service hello-node
service "hello-node" deleted
$ curl http://35.185.198.254:8080 --connect-timeout 10
curl: (28) Connection timed out after 10001 milliseconds
$ kubectl describe service hello-node
Error from server (NotFound): services "hello-node" not found
$ kubectl get pods
NAME                                READY    STATUS    RESTARTS   AGE
hello-node-952813049-2pf00         1/1      Running   0           9m
hello-node-952813049-dccn2         1/1      Running   0           9m
hello-node-952813049-zkxgf         1/1      Running   0           9m
$ kubectl delete deployment hello-node
deployment "hello-node" deleted
$ kubectl get pods
No resources found.
```

Google Kubernetes Engine Example Part 3 [cleanup]

1. Take down the service
2. Take down the replication controller for the pods

kubectl Setup

21

Copyright 2013-2018, RX-M LLC

- The kubectl command line tool is the generic front end for Kubernetes
 - Much like docker is to Docker
- Written in Go
 - Binaries for download on OS X, Linux and Windows
 - Build from Go source
 - Download as part of the Google Cloud SDK
 - Install with brew on OS X
- Reads config from ~/.kube/config
 - Configuration controlled by the config subcommand:
 - \$ kubectl config view
 - \$ kubectl config set-cluster
 - \$ kubectl config set-context

```
user@ubuntu:~$ wget https://storage.googleapis.com/kubernetes-release/release/$(curl -s
https://storage.googleapis.com/kubernetes-release/release/stable.txt)/bin/linux/amd64/kubectl
--2017-09-27 08:47:17-- https://storage.googleapis.com/kubernetes-release/release/v1.8.1/bin/linux/amd64/kubectl
Resolving storage.googleapis.com (storage.googleapis.com)... 172.217.5.112, 2607:f8b0:4005:808::2010
Connecting to storage.googleapis.com (storage.googleapis.com)|172.217.5.112|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 50359943 (48M) [application/octet-stream]
Saving to: 'kubectl'
```

```
kubectl          100%[=====>] 48.03M  297KB/s   in 1m 56s
2017-07-27 08:59:55 (266.8 KB/s) - 'kubectl' saved [50359943/50359943]
```

```
user@ubuntu:~$ chmod a+x kubectl
user@ubuntu:~$ ./kubectl get nodes
NAME        STATUS    AGE
127.0.0.1   Ready    48m
```

kubectl Config: Cluster and Context

- A kubectl **cluster** is a Kubernetes API endpoint for a given Kubernetes cluster
 - kubectl can be configured with multiple clusters
 - Only one context is active at a time and each context points to precisely one cluster
 - `kubectl config set-cluster` sets a cluster entry in kubeconfig [`~/.kube/config`]
 - `kubectl config set-cluster NAME [--server=server] [--api-version=apiversion] [--certificate-authority=path/to/certificate/authority] [--insecure-skip-tls-verify=true]`
 - Example
 - `$ kubectl config set-cluster e2e --server=https://15.22.31.4`
- A kubectl **context** defines the cluster connection metadata to be used with kubectl commands
- Set with: `kubectl config set-context`
 - Sets a context entry in kubeconfig [`~/.kube/config`]
 - Specifying a name that already exists will merge new fields on top of existing values for those fields
 - `kubectl config set-context NAME [--cluster=cluster_nickname] [--user=user_nickname] [--namespace=namespace]`
 - Example
 - `$ kubectl config set-context gce --user=cluster-admin`

```
user@ubuntu:~$ kubectl help config set-cluster
Sets a cluster entry in kubeconfig.
```

Specifying a name that already exists will merge new fields on top of existing values for those fields.

Examples:

```
# Set only the server field on the e2e cluster entry without touching other values. kubectl config set-cluster e2e --server=https://1.2.3.4
# Embed certificate authority data for the e2e cluster entry kubectl config set-cluster e2e --certificate-authority=~/.kube/e2e/kubernetes.ca.crt
# Disable cert checking for the dev cluster entry kubectl config set-cluster e2e --insecure-skip-tls-verify=true
```

Options:

```
--api-version='': api-version for the cluster entry in kubeconfig
--certificate-authority='': path to certificate-authority file for the cluster entry in kubeconfig
--embed-certs=false: embed-certs for the cluster entry in kubeconfig
--insecure-skip-tls-verify=false: insecure-skip-tls-verify for the cluster entry in kubeconfig
--server='': server for the cluster entry in kubeconfig
```

Usage:

```
kubectl config set-cluster NAME [--server=server] [--certificate-authority=path/to/certificate/authority][--insecure-skip-tls-verify=true] [options]
```

Use "kubectl options" for a list of global command-line options (applies to all commands).

kubectl Use

23

Copyright 2013-2018, RX-M LLC

- Syntax
 - `kubectl [command] [TYPE] [NAME] [flags]`
 - `command` – the operation to perform on the named resource(s)
 - `TYPE` – specifies the resource type
 - Specify singular, plural or abbreviated forms (deployment, deployments, deploy)
 - `NAME` – specifies the name of the resource
 - `flags` – optional flags
- Basic Resource Commands:
 - `get`
 - Display one or many resources
 - `describe`
 - Show details of a specific resource or group of resources
 - `create` (covered later)
 - Create a resource by filename or stdin
 - `delete`
 - Delete resources by filenames, stdin, resources and names, or by resources and label selector.
 - `logs`
 - Print the logs for a container in a pod
 - `attach`
 - Attach to a running container
 - `exec`
 - Execute a command in a container
- kubectl includes autocompletion support
 - Add it to your current shell
 - `source <(kubectl completion bash)`
 - To add to your profile so it is automatically loaded in future shells :
 - `echo "source <(kubectl completion bash)" >> ~/.bashrc`

Summary

- Kubernetes is a system for managing containerized applications across multiple hosts
- Kubernetes is itself a microservice based system and some or all of it may be executed via containers
- Kubernetes relies on etcd to maintain distributed watchable state
- Kubernetes can be deployed in a wide range of environments including bare metal OS, VM and cloud
- Kubernetes has been production ready for less than a year and is still changing and developing in significant ways

Lab 10

- Using Kubernetes with kubectl

11: Pods and Configs

Objectives

- Understand Kubernetes resource types
- Learn how to define resources with YAML files
- Define Pod and the operation of Pods on Kubernetes
- Understand the relationship between images, containers and Pods
- Start, Stop and monitor Pods

Kubernetes Concepts

28

Copyright 2013-2018, RX-M LLC

- **Cluster**
 - A cluster is a set of physical or virtual machines and other infrastructure resources used by Kubernetes to run your applications
- **Node**
 - A node is a physical or virtual machine running Kubernetes, onto which pods can be scheduled
- **Pod**
 - Pods are colocated groups of application containers
 - The smallest deployable units that can be created, scheduled, and managed with Kubernetes
 - Pods can be created individually, but it's recommended that you use a replication controller even if creating a single pod
- **Replica Set**
 - Replica Sets manage the lifecycle of pods
 - They ensure that a specified number of pods are running at any given time, by creating or killing pods as required
- **Deployment**
 - Allow initial deployment of pods through replica sets and rolling upgrades
- **Service**
 - Services provide a single, stable name and address for a set of pods
 - Act as basic load balancers
 - Services select the Pods to address using labels
- **Label**
 - Labels are key:value pairs used to organize and select groups of objects
- **Namespace**
 - A mechanism to partition resources created by users into a logically named group
 - Allows separate user communities to work in isolation

Kubernetes Process Control & Scheduling

29

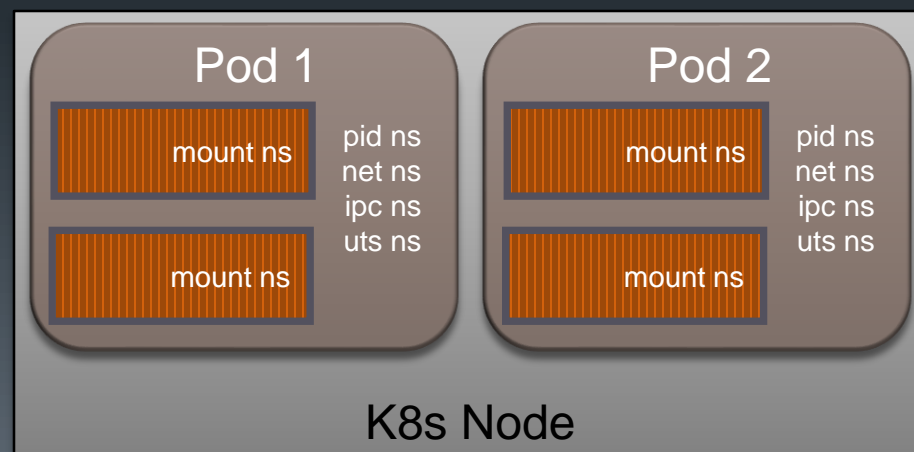
Copyright 2013-2018, RX-M LLC

- Desired State vs. Actual State
 - You **tell Kubernetes the state you desire** (run three copies of the nginx container)
 - **Spec**
 - Kubernetes works to bring the actual state inline with your desired state
 - **Status**
 - Replica Sets restart failed replicas, etc.
 - This is the eventual consistency model at work at a higher level of abstraction
- The Kubernetes scheduler spreads pods across the cluster and uses different nodes for matching pod replicas when possible
- Many ancillary Kubernetes services may be deployed, usually themselves in containers
 - The default deployment model for the kubelet will likely be changed to in-container in the near future
 - Docker installation example from section 2 uses containerized kubelet

Kubernetes Goal:
Make
Status = Spec

Pods

- Many times a set of containers run as a team on the same host
 - A data loader with a data server
 - A log compressor/saver process with a log server
- These containers usually need to be located together
 - You want to ensure that they do not become separated during dynamic placement
- In Kubernetes **a pod is a set of containers that are placed and scheduled together as a unit** on a Kubernetes Node
 - Kubernetes Node is a worker computer in the cluster
- By working to place a group of pods, Kubernetes can pack lots of work onto a node in a reliable way
- In Kubernetes, Pods are defined using YAML (or, the more recently supported JSON) files



Config File Commonalities

31

Copyright 2013-2018, RX-M LLC

- YAML formatted
- Main parts
 - **# Comment**
 - **apiVersion**
 - Kubernetes API version to build with
 - Presently should always be v1
 - **kind**
 - Kubernetes object type
 - Pod, Deployment, Service, Namespace, etc...
 - **metadata**
 - Name, labels and other data describing the object
 - **spec**
 - Instructions for building the object
 - The desired state
- Construction
 - **kubectl create -f filename**

```
# mypod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: web-pod
spec:
  containers:
  - name: web-container
    image: bitnami/apache:latest
    ports:
    - containerPort: 80
```

Pod Definition

32

Copyright 2013-2018, RX-M LLC

- The simplest pod definition describes the deployment of a single container
- For example, an nginx web server pod might be defined as such:

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx
    ports:
    - containerPort: 80
```

- A pod definition is a declaration of a **desired state**
- Desired state is a very important concept in the Kubernetes model
- Many things present a desired state to the system, and it is Kubernetes' responsibility to make sure that the current state matches the desired state
 - For example, when you create a Pod, you declare that you want the containers in it to be running
 - If the containers happen to not be running (e.g. program failure, ...), Kubernetes will continue to (re-)create them for you in order to drive them to the desired state
 - **Spec** describes build time (what we want), and **Status** describes runtime (what we have).
 - This process continues until the Pod is deleted

Pod Walkthrough

Part 1

```
user@ubuntu:~/pods$ vim web.yaml
user@ubuntu:~/pods$ cat web.yaml
apiVersion: v1
kind: Pod
metadata:
```

```
  name: web-pod
spec:
  containers:
  - name: web-container
    image: bitnami/apache:latest
    ports:
    - containerPort: 80
```

```
user@ubuntu:~/pods$ kubectl create -f web.yaml
```

```
pod "web-pod" created
```

```
user@ubuntu:~/pods$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
web-pod	0/1	Pending	0	13s

```
user@ubuntu:~/pods$ kubectl describe pod web-pod
```

```
Name:          web-pod
Namespace:     default
Image(s):      bitnami/apache:latest
Node:          127.0.0.1/127.0.0.1
Start Time:    Fri, 20 Jan 2017 06:29:30 -0800
Labels:        <none>
Status:        Running
IP:            172.17.0.2
Controllers:   <none>
Containers:
  web-container:
    Container ID:   docker://76b4e62aea4a1e72ed7c26752194541453d8f4e9d99e681b86ef5f2ad969412b
    Image:          bitnami/apache:latest
    Image ID:       docker-pullable://bitnami/apache@sha256:4b7dc6a2b0e06c787bc6f38472926b270cb97b14fd1e47637a469ab8a56f9682
    Port:           80/TCP
    State:          Running
      Started:      Fri, 20 Jan 2017 06:29:55 -0800
    Ready:          True
    Restart Count:   0
    Volume Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-c3m82 (ro)
    Environment Variables:  <none>
Conditions:
  Type           Status
  Initialized     True
  Ready          True
  PodScheduled    True
...
```

Pod Walkthrough

Part 2

```
user@ubuntu:~/pods$ kubectl describe pod web-pod
```

...

```
Volumes:
  default-token-c3m82:
    Type:          Secret (a volume populated by a Secret)
    SecretName:     default-token-c3m82
```

```
QoS Class:   BestEffort
```

```
Tolerations: <none>
```

```
Events:
```

FirstSeen	LastSeen	Count	From	SubObjectPath	Type	Reason	Message
13m	13m	1	{default-scheduler }		Normal	Scheduled	Successfully assigned web-pod to 127.0.0.1
13m	13m	1	{kubelet 127.0.0.1}	spec.containers{web-container}	Normal	Pulling	pulling image "bitnami/apache:latest"
13m	13m	1	{kubelet 127.0.0.1}	spec.containers{web-container}	Normal	Pulled	Successfully pulled image "bitnami/apache:latest"
13m	13m	1	{kubelet 127.0.0.1}	spec.containers{web-container}	Normal	Created	Created container with docker id 76b4e62aea4a; Security:[seccomp=unconfined]
13m	13m	1	{kubelet 127.0.0.1}	spec.containers{web-container}	Normal	Started	Started container with docker id 76b4e62aea4a

```
user@ubuntu:~/pods$ curl -s http://172.17.0.2:80
<html><head/><body><h1>It works!</h1></body></html>
```

```
user@ubuntu:~/pods$ kubectl exec web-pod -- ps -ef
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
root	1	0	0	01:10	?	00:00:00	tini -- nami start --foreground apache
root	36	1	0	01:10	?	00:00:01	/opt/bitnami/nami/runtime/node --max_semi_space_size=150 /opt/bitnami/nami/index.js start --foreground apache
root	49	1	0	01:10	?	00:00:00	/opt/bitnami/apache/bin/httpd -f /opt/bitnami/apache/conf/httpd.conf
daemon	50	49	0	01:10	?	00:00:00	/opt/bitnami/apache/bin/httpd -f /opt/bitnami/apache/conf/httpd.conf
daemon	51	49	0	01:10	?	00:00:00	/opt/bitnami/apache/bin/httpd -f /opt/bitnami/apache/conf/httpd.conf
daemon	52	49	0	01:10	?	00:00:00	/opt/bitnami/apache/bin/httpd -f /opt/bitnami/apache/conf/httpd.conf
daemon	53	49	0	01:10	?	00:00:00	/opt/bitnami/apache/bin/httpd -f /opt/bitnami/apache/conf/httpd.conf
daemon	54	49	0	01:10	?	00:00:00	/opt/bitnami/apache/bin/httpd -f /opt/bitnami/apache/conf/httpd.conf
root	69	0	0	01:40	?	00:00:00	ps -ef

```
user@ubuntu:~/pods$ kubectl delete pod web-pod
pod "web-pod" deleted
```

```
user@ubuntu:~/pods$ kubectl get pods
No resources found.
```

```
user@ubuntu:~/pods$
```

What is YAML

35

Copyright 2013-2018, RX-M LLC

- Rhymes with camel
- Stands for YAML Ain't Markup Language
 - Was originally Yet Another Markup Language but the new reverse backronym and antiestablishment attitude is apparently cooler
- Human-readable data serialization format that takes concepts from programming languages and ideas from XML and electronic mail (RFC 2822)
- Designed to be easily mapped to data types common to most high-level languages:
 - List [array/set]
 - Associative array [hash/object/map]
 - Scalar
- Friendly to ad hoc grep/Python/Perl/Ruby operations
- Eschews enclosures (quotation marks, brackets, braces, open/close-tags, etc.)
- Data structure hierarchy is maintained by outline indentation
 - Spaces only, no tabs!



Docker uses YAML in some human targeted file areas but most files are JSON formatted

```
---
receipt:  VMW Invoice
date:     2012-08-06
customer:
  given:   Dorothy
  family:  Gale
```

- YAML offers an indented and an in-line styles

- Lists

- Conventional block format uses a hyphen+space to begin a new item in list.

```
--- # Favorite movies
- Casablanca
- North by Northwest
- The Man Who Wasn't There
```

- Optional inline format is delimited by comma+space and enclosed in brackets (similar to JSON)

```
--- # Shopping list
[milk, pumpkin pie, eggs, juice]
```

- Associative arrays

- Keys are separated from values by a colon+space. Indented Blocks use new lines to separate key: value pairs; Inline Blocks use comma+space to separate the key: value pairs between braces

```
--- # Indented Block
  name: John Smith
  age: 33
--- # Inline Block
{name: John Smith, age: 33}
```

- Block literals

- Strings do not require quotation
- Newlines preserved (By default first line indent and trailing space is stripped though other behavior can be specified)

```
--- |
  There once was a man from Ealing
  Who got on a bus to Darjeeling
  It said on the door
  "Please don't spit on the floor"
  So he carefully spat on the ceiling
```

- Newlines folded (Folded text converts newlines to spaces and removes leading whitespace)

```
--- >
  Wrapped text
  will be folded
  into a single
  paragraph

  Blank lines denote
  paragraph breaks
```

- Hierarchical combinations of elements

- Lists of associative arrays

```
- {name: John Smith, age: 33}
- name: Mary Smith
  age: 27
```

- Associative arrays of lists

```
men: [John Smith, Bill Jones]
women:
  - Mary Smith
  - Susan Williams
```

Pod Configuration Files

37

Copyright 2013-2018, RX-M LLC

- Pods can be configured using YAML or JSON configuration files
 - Much like Docker Compose
 - `kubectl create -f mynewpod.yml`
- Fields for Pod configuration include:
 - **apiVersion**: Currently v1
 - **kind**: Always Pod
 - **metadata**: An object containing:
 - **name**: The name of this pod (required if generateName is not specified)
 - **labels**: Optional arbitrary key:value pairs used for grouping and targeting by other resources and services
 - **spec**: The pod specification
 - **volumes[]**: List of volumes that can be mounted by containers in the pod
 - **restartPolicy**: applies to all containers in the pod
 - **nodeSelector**: node label required for pod to run on a given node
 - **nodeName**: specific node to schedule pod on
 - **terminationGracePeriodSeconds**: number of seconds node is given to shutdown before kill
 - **hostNetwork**: run pod in host net namespace
 - **hostPID**: run pod in host PID namespace
 - **hostIPC**: run pod in host IPC namespace
 - **imagePullSecrets**: secrets to use for pulling images (e.g. DockerConfig for private reg access)
 - **containers[]**: A list of containers belonging to the pod
 - **name**: Name of the container
 - **image**: Docker image name
 - **command[]**: command line (like Docker ENTRYPOINT)
 - **args[]**: entry point arguments (like Docker CMD)
 - **env[]**: env vars
 - **ports[]**: port mappings
 - **containerPort**: port to expose
 - **protocol**: port protocol
 - **hostIP**: host IP to bind to
 - **hostPort**: host port to bind to
 - **volumes[]**: volumes which can be mounted
 - **restartPolicy**: Always, OnFailure, Never
 - **workingDir**: cwd for container proc

```
user@ubuntu:~/pods$ cat hello.yaml
apiVersion: v1
kind: Pod
metadata:
  name: hello-world
spec: # specification of the pod's contents
  restartPolicy: Never
  containers:
  - name: hello
    image: "ubuntu:14.04"
    env:
    - name: MESSAGE
      value: "hello world"
    command: ["/bin/sh", "-c"]
    args: ["/bin/echo \"${MESSAGE}\""]
```

```
user@ubuntu:~/pods$ kubectl create -f hello.yaml
pod "hello-world" created
user@ubuntu:~/pods$ kubectl get pod
NAME          READY    STATUS    RESTARTS   AGE
hello-world   0/1      Pending   0           15s
user@ubuntu:~/pods$ kubectl logs hello-world
hello world
user@ubuntu:~/pods$ kubectl get pod
user@ubuntu:~/pods$
```

Pod Spec Reference:

https://kubernetes.io/docs/api-reference/v1/definitions/#_v1_podspec

Container Spec Reference:

https://kubernetes.io/docs/api-reference/v1/definitions/#_v1_container

Resource Limits

- You can specify the desired CPU and memory for containers in specs to ensure the scheduler chooses a node with the appropriate resources
- CPU and memory are each a resource type
 - CPU is specified in units of cores
 - `spec.container[].resources.requests.cpu`
 - Memory is specified in units of bytes
 - `spec.container[].resources.requests.memory`
- Each container of a Pod can optionally specify **constraints**
 - `spec.container[].resources.limits.cpu`
 - `spec.container[].resources.limits.memory`
- Default values are cluster configured
 - If value of requests is not specified, they are set to be equal to limits by default
 - Resource limits must be greater than or equal to resource requests
- requests/limits can only be specified on individual containers
 - Pod request/limits are simply the sum of their container request/limits
- If the scheduler cannot find any node where a pod can fit, then the pod will remain unscheduled until a place can be found

```
apiVersion: v1
kind: Pod
metadata:
  name: frontend
spec:
  containers:
    - name: db
      image: mysql
      resources:
        requests:
          memory: "64Mi"
          cpu: "250m"
        limits:
          memory: "128mi"
          cpu: "500m"
    - name: wp
      image: wordpress
      resources:
        requests:
          memory: "64Mi"
          cpu: "250m"
        limits:
          memory: "128mi"
          cpu: "500m"
```

Pod Phases

39

Copyright 2013-2018, RX-M LLC

- Phases
 - Describe the macro states in the lifecycle of a Kubernetes resource
 - Not a comprehensive state machine
- Pod Phase
 - **Pending**
 - The pod has been accepted by the system, but one or more of the container images has not been created
 - Includes time before being scheduled as well as time spent downloading images over the network
 - **Running**
 - The pod has been bound to a node, and all of the containers have been created
 - At least one container is still running, or is in the process of starting or restarting
 - **Succeeded**
 - All containers in the pod have terminated in success, and will not be restarted
 - **Failed**
 - All containers in the pod have terminated, at least one container has terminated in failure (exited with non-zero exit status or was terminated by the system)
 - **Unknown**
 - For some reason the state of the pod could not be obtained, typically due to an error in communicating with the host of the pod

- The Pause container is often referred to as the pod infrastructure container and is used to set up and hold the networking namespace and resource limits for each pod

```
user@ubuntu:~/pods$ cat long.yaml
apiVersion: v1
kind: Pod
metadata:
  name: long-running
spec: # specification of the pod's contents
  containers:
  - name: long
    image: "ubuntu:14.04"
    command: ["/usr/bin/tail", "-f", "/dev/null"]
```

```
user@ubuntu:~/pods$ kubectl create -f long.yaml
pod "long-running" created
```

```
user@ubuntu:~/pods$ kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
long-running	0/1	Running	0	3s

```
user@ubuntu:~/pods$ sudo docker container ls -f "name=long-running"
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	NAMES
8fe590c80a12	ubuntu:14.04	"/usr/bin/tail -f /de"	2m ago	Up 2m	k8s_long.94251326_long...
8bf3dbb786cc	gcr.io/google_containers/pause-amd64:3.0	"/ pause "	2m ago	Up 2m	k8s_POD.6d00e006_long...

Pod Patterns

41

Copyright 2013-2018, RX-M LLC

■ Patterns

■ Sidecar

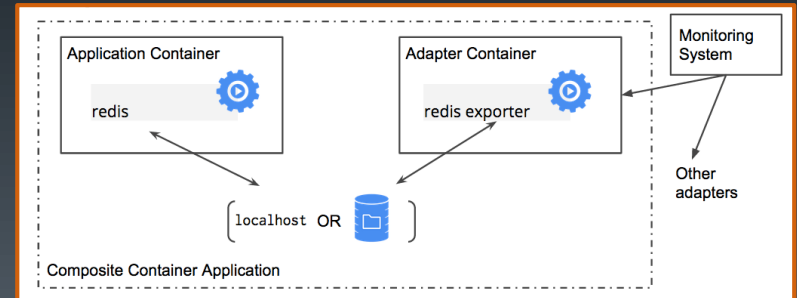


- Sidecars extend and enhance the "main" container in the Pod
- Example: Nginx web server container; add a container that syncs the file system with a git repository, share the file system between the containers and you have built Git push-to-deploy

■ Ambassador

- Ambassadors proxy a Pod local connection to the world outside
- Example: Redis cluster with read-replicas and a single write master; create a Pod that groups the main application with a Redis ambassador container which splits reads and writes, sending them on to the appropriate servers

■ Adapter

- Adapter containers standardize and normalize output
- Example: A task monitoring N different applications where each application has a different way of exporting monitoring data (e.g. JMX, StatsD, application specific statistics) but every monitoring system expects a consistent and uniform data model for the monitoring data it collects



- Brendan Burns, Distinguished Engineer at Microsoft and former Software Engineer at Google
<https://research.google.com/pubs/pub45406.html>

Summary

- Kubernetes resource types provide the basic building blocks for Kubernetes based applications
- YAML files are used to specify resources and the metadata associated with them
- Pod are the unit of application deployment in Kubernetes
- Pods are atomic and scheduled together on a single node as a unit
- A Pod may contain one or more user defined containers
- A Pod has an infrastructure container (Pause) used as a place holder for the Pods namespaces and CGroups
- Pods share network and IPC namespaces
- kubectl offers several commands to work with pods
 - create
 - delete
 - describe
 - get
 - logs
- A container engine is used behind the scenes to execute and manage containers (e.g. Docker)
- The container engine is accessible directly for debugging and certain management chores

Lab 11

- Pods and Config Files